

# A Preliminary Evaluation on Many-to-one Virtualization

**Jin Zhang**, Zhuocheng Ding, Yubin Chen,  
Zhengwei Qi, Haibing Guan



上海交通大學

SHANGHAI JIAO TONG UNIVERSITY

**ACM TURC 2019**

This work is supported by the National Key Research & Development Program of China 2016YFB1000502

# CONTENTS


- 1 Motivation and Background
- 2 GiantVM: Design
- 3 Evaluation and Illustration

# CONTENTS

- 1 Motivation and Background

---
  - 2 GiantVM: Design

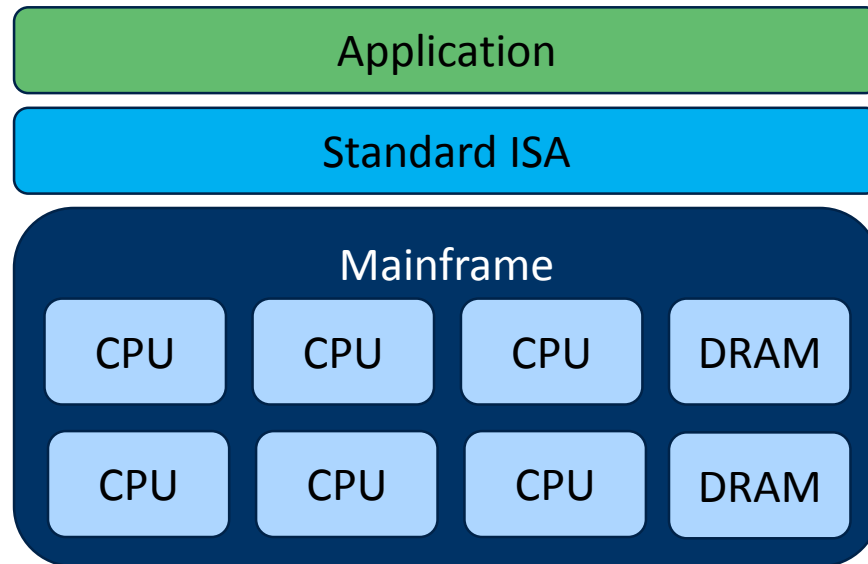
---
  - 3 Evaluation and Illustration

---
- 

# Resource Aggregation



- How can you aggregate resources in case of the end of Moore's Law?



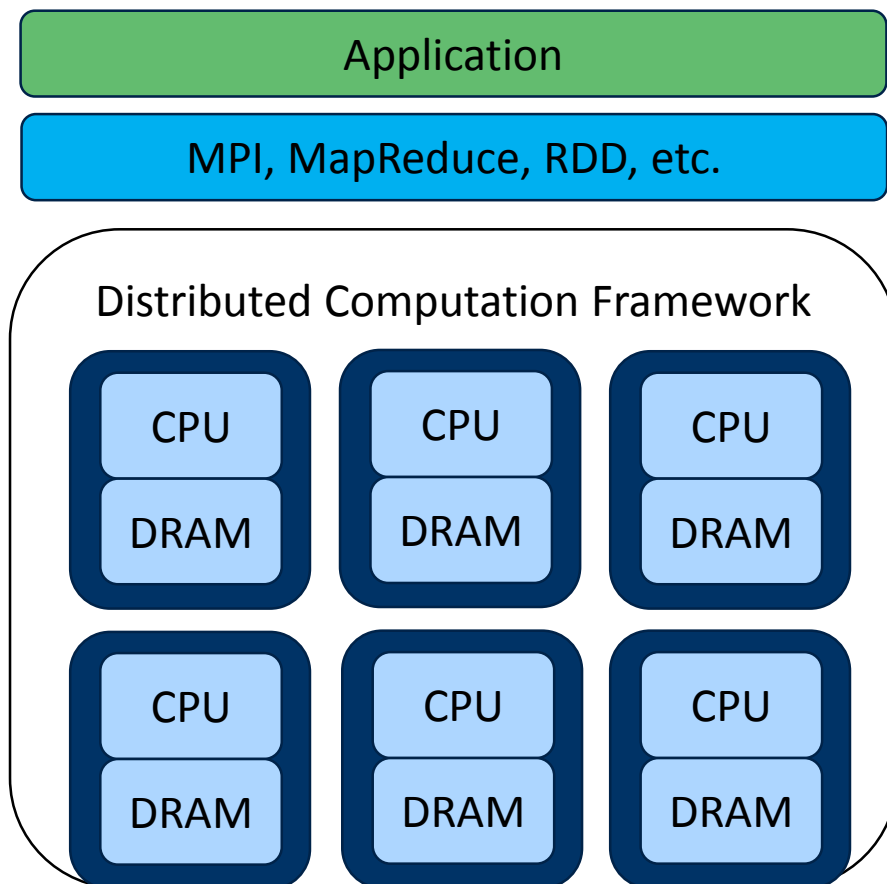
Scale-Up:  
one machine with many CPUs, memory, etc.

- + No need to port existing software
- Expensive

# Resource Aggregation



- How can you aggregate resources in case of the end of Moore's Law?



Scale-Out:  
many machines with CPUs, memory, etc.

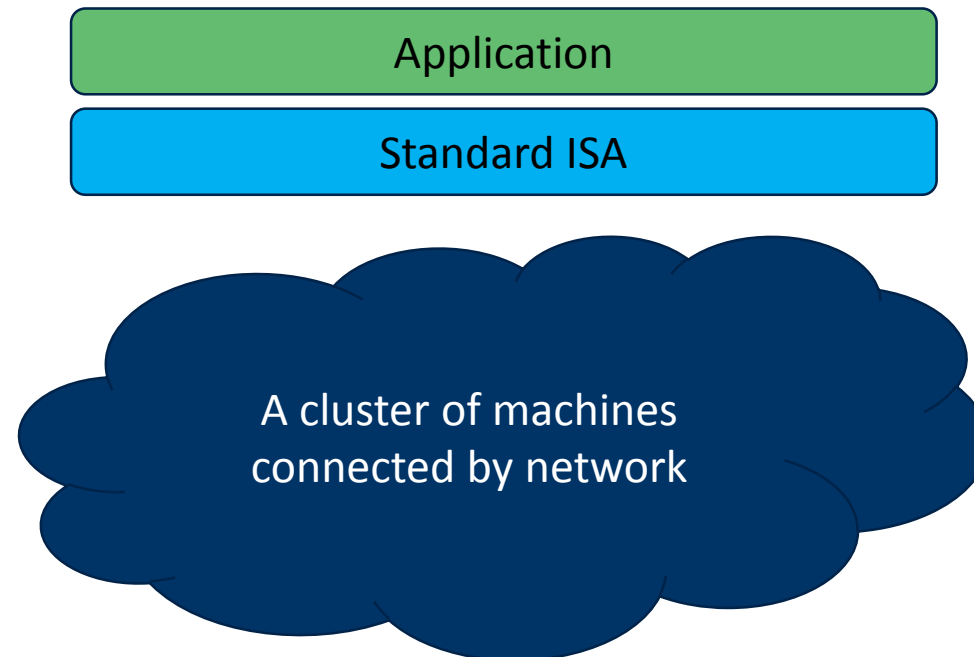
- + Affordable
- A huge engineering effort to port existing software

# Single System Image



- **Single System Image (SSI)**

- A single OS instance running on multiple machines.
- No need to modify existing software.




# Single System Image



- Recently engaged by the high-speed network, such as RDMA, NVMe over Fabrics, Intel OmniPath, etc.
- Has different types.

Memory aggregation




Infiniswap (NSDI'17),  
Remote Regions (ATC'18),  
Memory Blade (ISCA'09, HPCA'12)

# Single System Image



- Recently engaged by high-speed network, such as RDMA, NVMe over Fabrics, Intel OmniPath, etc.
- Has different types.

I/O devices  
aggregation,  
including  
heterogeneous ones.



NFS, iSCSI, ReFlex (ASPLOS'17),  
Decibel (NSDI'17),  
PolarFS (PVLDB'18)




# Single System Image



- Recently engaged by high-speed network, such as RDMA, NVMe over Fabrics, Intel OmniPath, etc.
- Has different types.

## Distributed OS




HeliOS (SOSP'09), fos (SIGOPS'09),  
Popcorn (EuroSys'15),  
LegoOS (OSDI'18)

# Single System Image



- Recently engaged by high-speed network, such as RDMA, NVMe over Fabrics, Intel OmniPath, etc.
- Has different types.




ScaleMP, TidalScale,  
vNUMA (ATC'09),  
GiantVM (Our work)


# Rethink System Virtualization



One-to-many virtualization  
(Classical virtualization)



Our approach:  
Many-to-one virtualization




Implement SSI by many-to-one virtualization

# CONTENTS

- 1 Motivation and Background
- 2 **GiantVM: Design**
- 3 Evaluation and Illustration

# GiantVM: Overview of Architecture



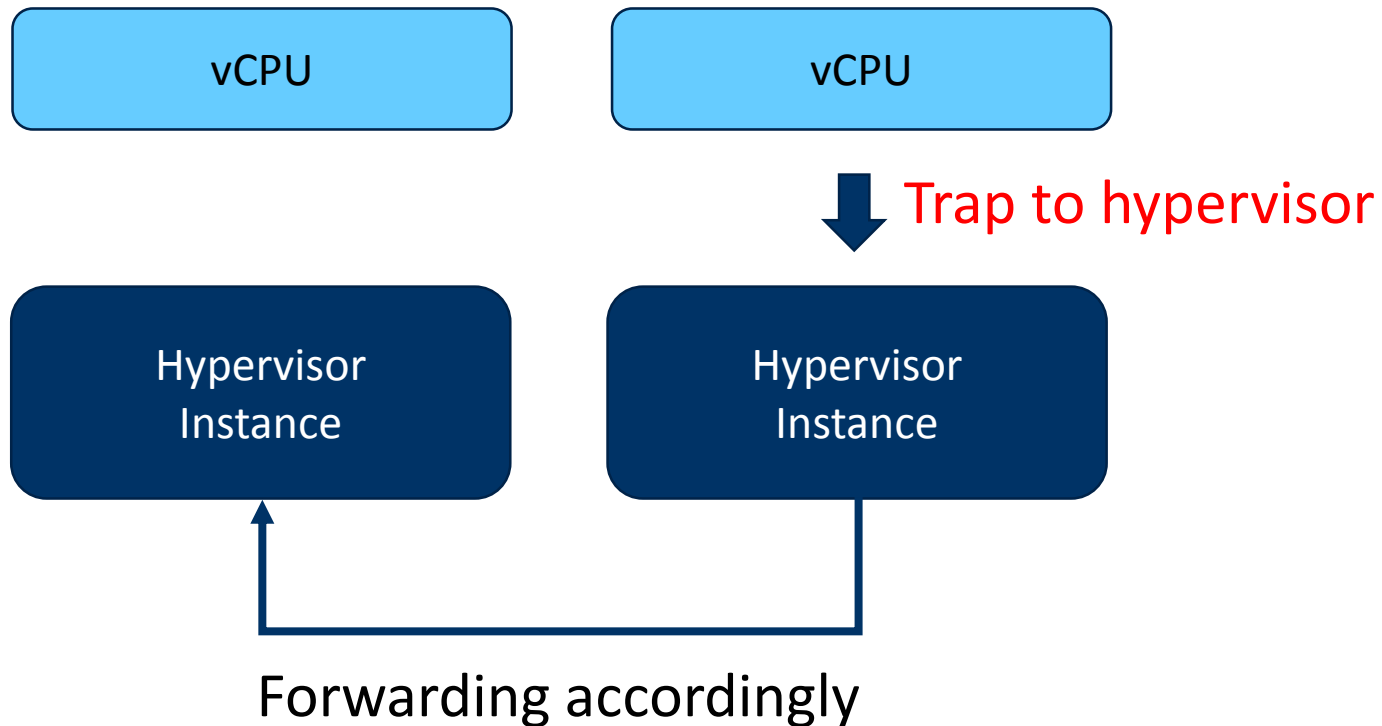
- Each node runs with a hypervisor instance.
- Each instance is in charge of partial resources.
- The global consistent view seen by the guest is maintained by the instances.

# GiantVM: Design



- **The Von Neumann architecture consists of three components:**
  - CPU
  - I/O
  - Memory
- **And GiantVM should make these components distributed.**
- **CPU is distributed originally.**
- **I/O is centralized and usually multiplexed by device drivers.**
- **Distributed memory is also known as Distributed Shared Memory (DSM).**

# GiantVM: CPU and I/O Virtualization



- **Only interactions between different nodes should be considered.**
  - E.g., inter-processor interrupts (IPI), memory-mapped I/O (MMIO), port I/O (PIO), etc.
- **These instructions are simulated by the hypervisor originally. GiantVM intercepts them by forwarding instructions to the proper remote node.**



# GiantVM: Memory Virtualization

- **A straight-forward Ivy (K. Li, 1989) implementation on the EPT.**
  - Although a plethora of troubles when porting to QEMU-KVM...
- **Although Ivy only implements sequential consistency (an ancient and somewhat low-efficient memory model), there is no alternatives.**
  - Our target platform x86 follows x86-TSO.
  - In x86-TSO, write can be delayed until a 'mfence' is executed.
  - However, we have no way to capture execution of 'mfence'. Thus we can only apply a conservative strategy.





# Ivy Protocol in a Nutshell

- Each page of memory space has one of three states:

Privilege	State
Read and write	<b>M</b> odified
Read only	<b>S</b> hared
Cannot read or write	<b>I</b> nvalid

- Some terminologies:

Term	Description
Owner	The node holding the latest data.
Copyset	The nodes holding the valid data (owner is always included).

# Ivy Protocol in a Nutshell



- Two operations (read and write) on three states:

	read	write
Modified	✓	✓
Shared	✓	Invalid other copies. Change to the owner, <i>Modified</i>
Invalid	Ask owner for latest page. Change to the <i>Shared</i>	Ask owner for latest page Invalid other pages. Change to the owner, <i>Modified</i>

# CONTENTS

- 1 Motivation and Background
- 2 GiantVM: Design
- 3 Evaluation and Illustration

# Experiment Setup



Cluster	
Number of Machines	4
CPU	8-core Intel Xeon E5-2620 v4
DRAM	128GB
Ethernet NIC	Broadcom NetXtreme BCM5720 Gigabit Ethernet
RDMA HCA	ConnectX-3 MCX354A-FCBT 56Gbps InfiniBand
OFED Version	Mellanox OFED v2.2-1
Disk	SEAGATE ST9300605SS
OS	Ubuntu 16.04
Control Group	
CPU	24-core Intel Xeon E5-2650 v4
DRAM	64GB
OS	Ubuntu 16.04
Guest Configuration	
DRAM	64GB
OS	Ubuntu 16.04


Baseline

# Microbenchmarks



- **Sysbench-memory**

- Best scenarios: read-only operations or write local pages.



- **Note the similarity between MESI or MOESI protocols used in cache system and Ivy protocol used in the DSM.**

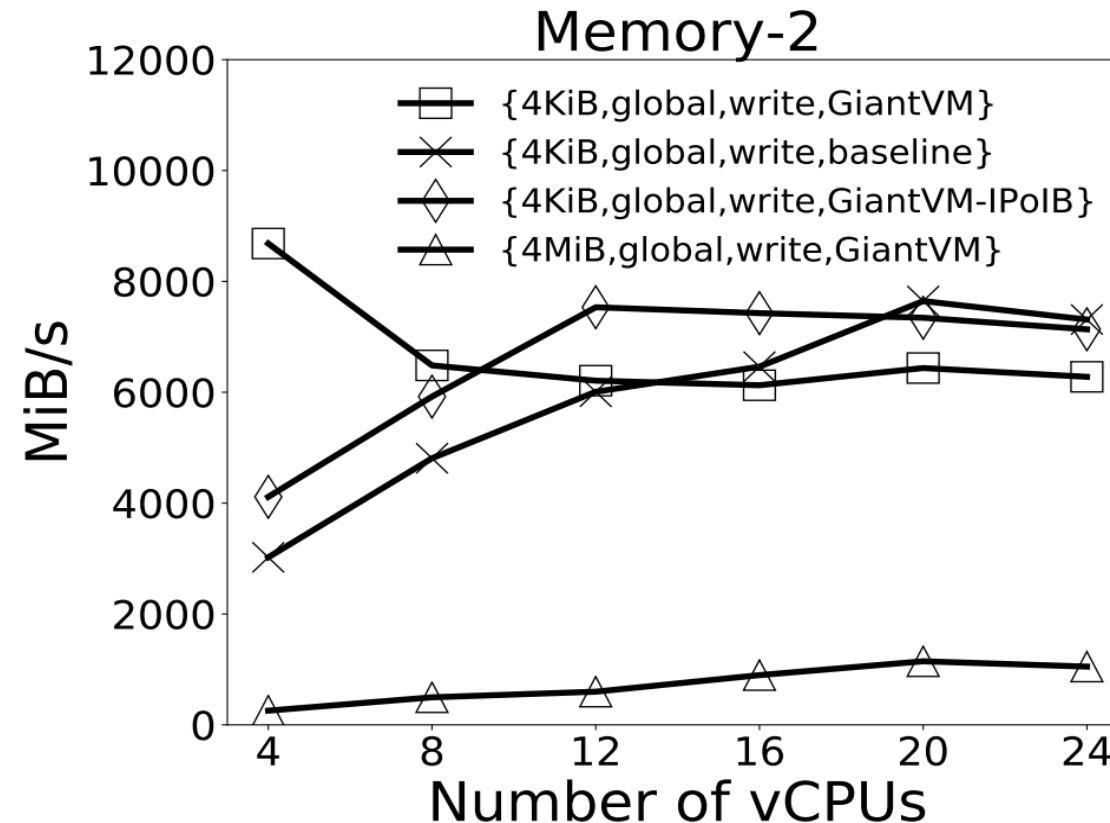
- Read-only or write-local are the optimum which leads no state transmissions.

# Microbenchmarks



- **Sysbench-memory**

- Worst scenarios in theory: write global shared page.
- However, there is an interesting “slower is faster” phenomenon.

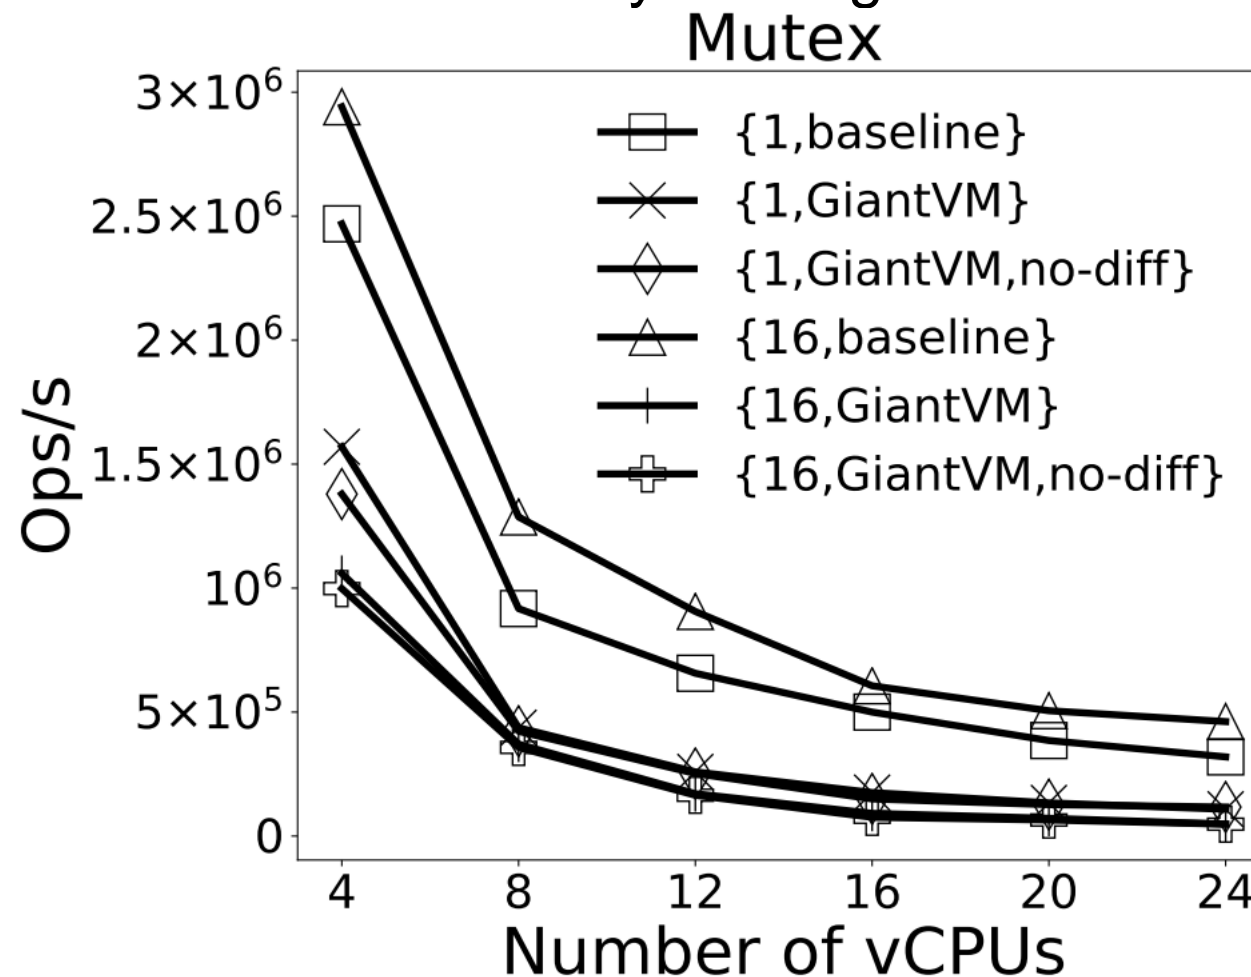


# Microbenchmarks



- **Sysbench-mutex**


- Another worst scenarios in theory: write global shared mutex(es).





# Illustration

- **A *shared unit*** means a cacheline in the cache system or a page in the DSM system.
- **Thrashing**
  - One shared unit is written by multiple vCPUs. The shared unit therefore ping-pongs across different vCPUs.
- **False Sharing**
  - Multiple memory objects are allocated in one shared unit.
  - The access of one object of those brings another one together.
  - Sol. Allocate cacheline (page size) alignment memory objects.





# Illustration




- **Now back to write-global scenario.**
  - Sysbench-Memory and Sysbench-Mutex are different.
  - Memory is *throughput-oriented*, when the local vCPU **can** make progress if remote Invalid message has not arrived yet.
  - Mutex is *latency-oriented*, when the local vCPU **cannot** make progress if remote Invalid message has not arrived yet.

# Illustration



- **Three scenarios**

- 1. A slow Invalid message.
- 2. A fast Invalid message.
- 3. The working set spans across the multiple shared units.




## Legend



Invalid State



Modified State



Processing




Invalid Message

# Illustration



- **Three scenarios**

- 1. A slow Invalid message.
- **2. A fast Invalid message.**
- 3. The working set spans across the multiple shared units.



## Legend



Invalid State



Modified State



Processing




Invalid Message

# Illustration



- **Three scenarios**

- 1. A slow Invalid message.
- 2. A fast Invalid message.
- **3. The working set spans across the multiple shared units.**

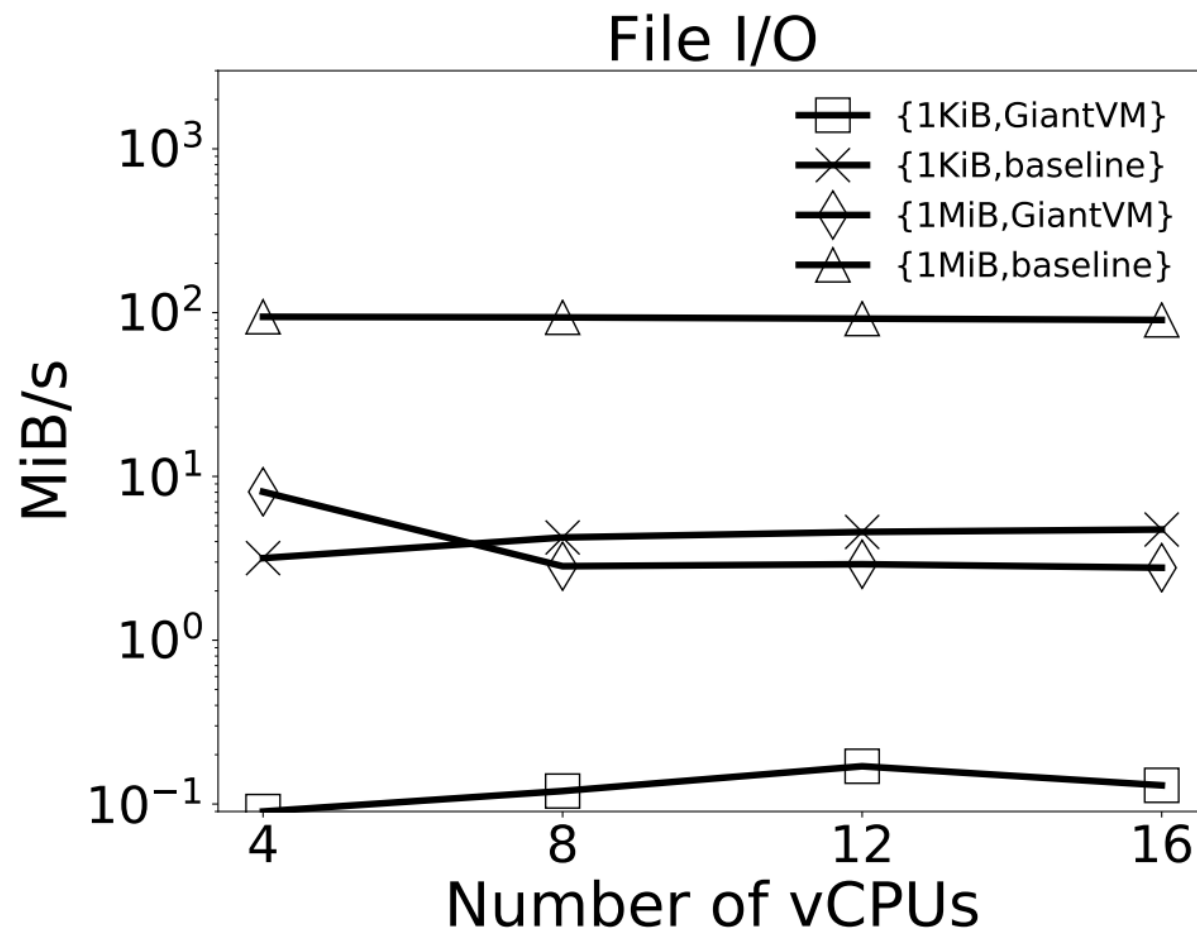


# Microbenchmark



- **Sysbench-File I/O**

- A catastrophic consequence of File I/O with **20x** slow down.



# Microbenchmark



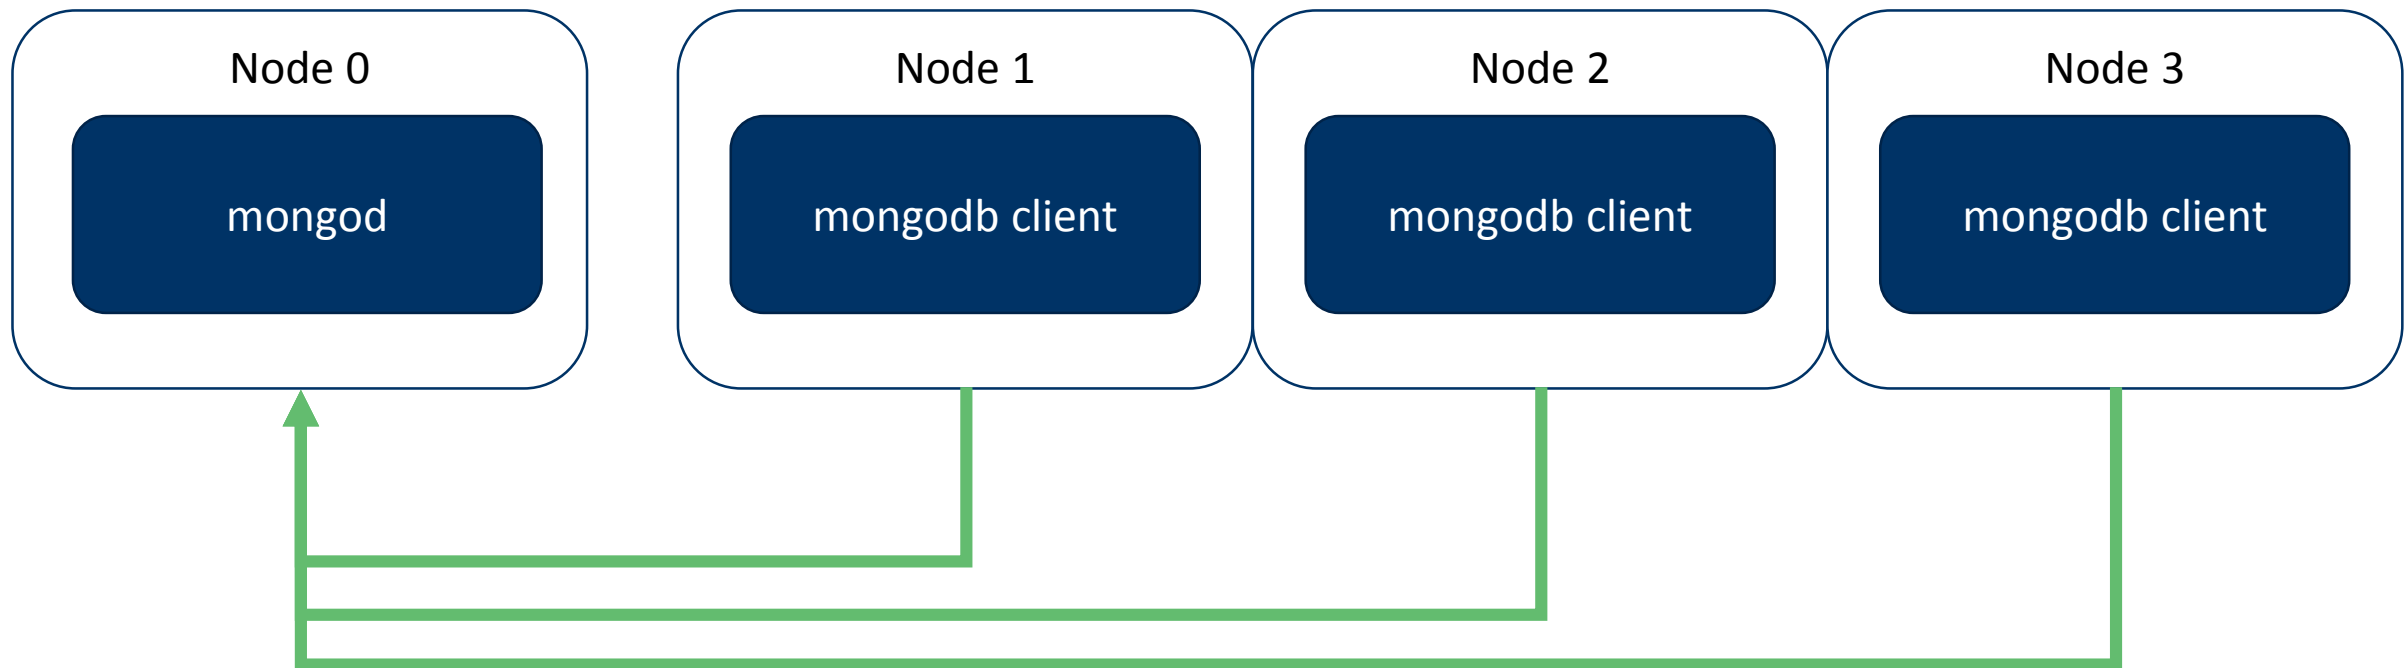
- **Direct usage of applications causes negative consequence.**
  - The I/O-intensive workload frequently traps to kernel and kernel pages are shared across all the vCPUs.
  - The shared pages lead to thrashing and false sharing.

**Monolithic kernel like Linux may not be a good idea!**

# Application Evaluation: Data Colocation



- **Some ideas borrowed from *Multikernel*:**
  - The applications as well as device drivers are binding to specific cores.
  - The communication between different cores is via shared-memory RPC.
- **Consider an example of MongoDB:**

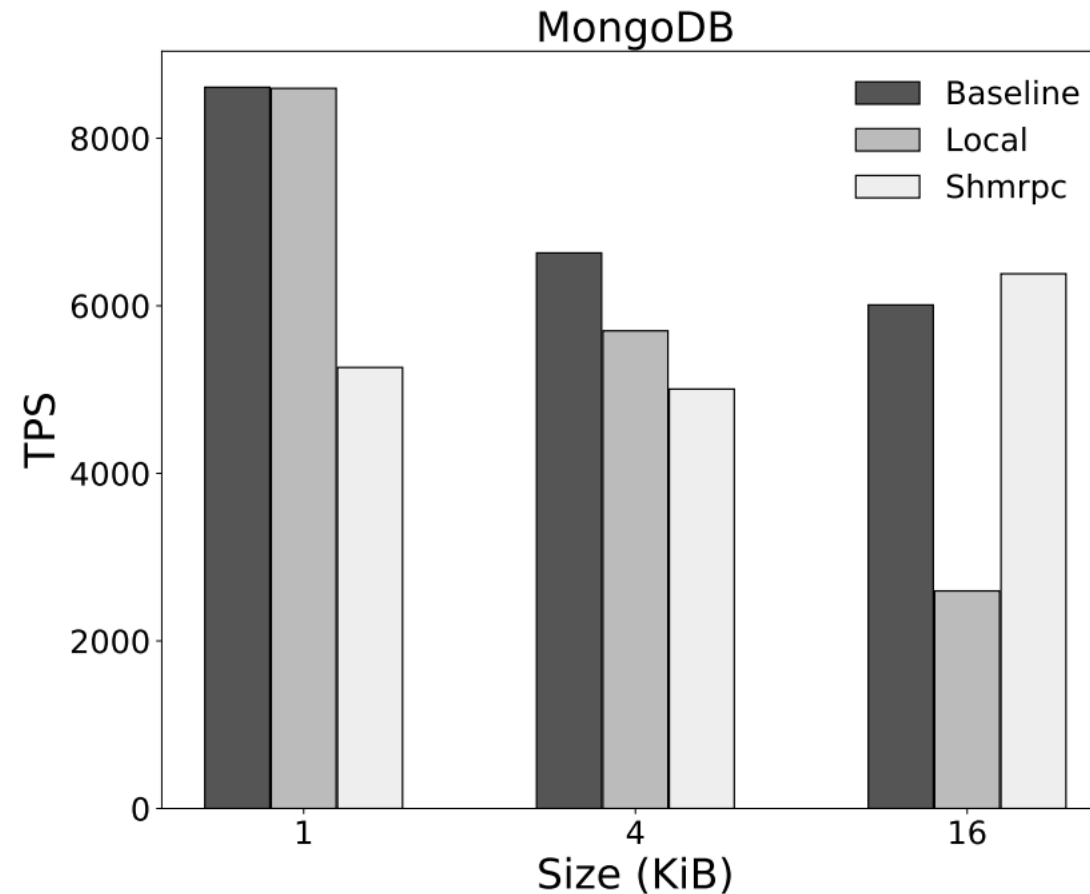


Shared Memory RPC

# Application Evaluation: Data Colocation



- **MongoDB spins on inserting/removing randomly generated documents.**

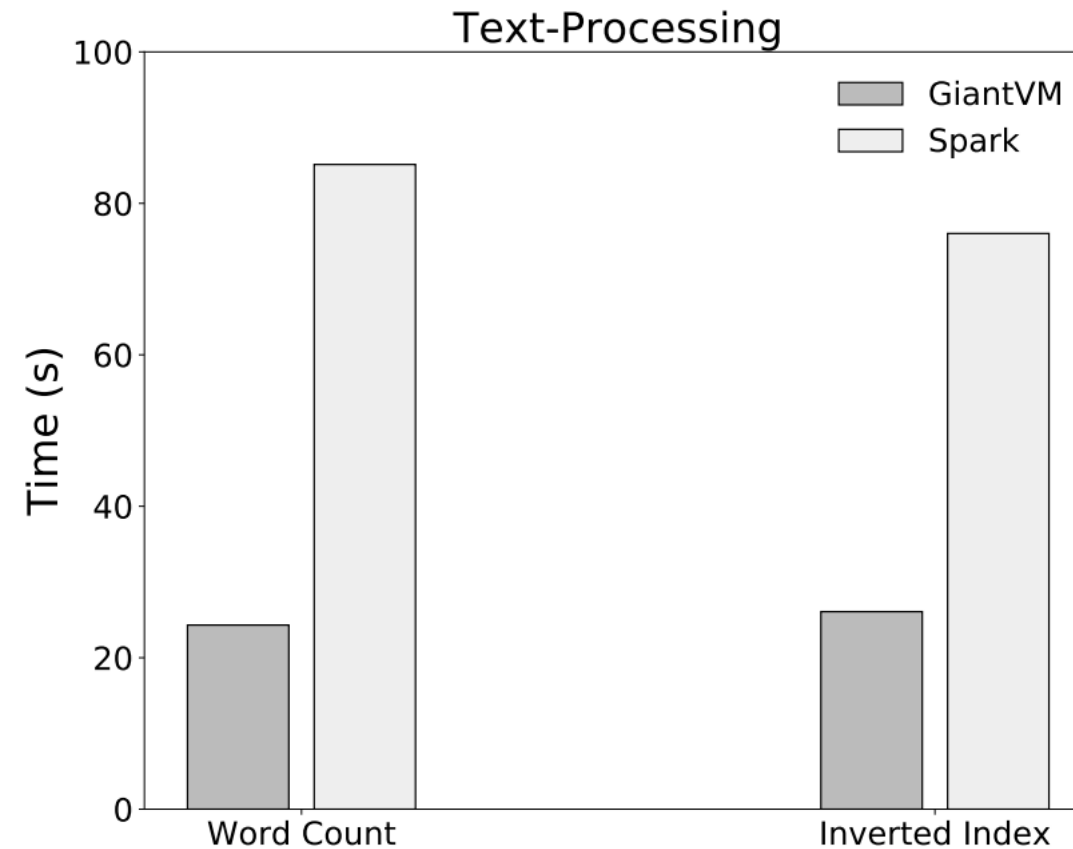






# Application Evaluation: Spark

- Comparison of two text-processing programs with Spark:



# Future Work



- **The integration with multikernel (Barrelfish) is work in progress.**
  - Multikernel brings additional benefits, especially fault-tolerance.
- **Relaxing memory model.**
  - ‘mfence’ instructions can be simulated by para-virtualization.
  - Only in the kernel space.

# Conclusion



- **GiantVM verifies the viability of applying many-to-one virtualization to the SSI.**
- **There are some interesting features when running specific workloads in GiantVM.**
- **A monolithic kernel like Linux may not be the ideal architecture. Instead, ideas from multikernel brings the enhancement of the performance.**

# Q&A

Our website: <http://tcloud.sjtu.edu.cn/>



上海交通大學  
SHANGHAI JIAO TONG UNIVERSITY